



**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich



# Introduction to GPUs in HPC

Ben Cumming, CSCS  
July 19, 2017



**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich

# Using GPUs in Your Application

---

Rule #1: **don't** develop your own GPU code!

# Libraries

There are many open libraries for GPUs.

- **cuBLAS**: Dense linear algebra primitives.
- **Thrust**: C++ STL-like algorithms and containers.
- **cuRAND** and **Random123**: Random numbers.
- **cuFFT**: FFT
- **Kokkos**: Generic performance portable parallel motifs.

... And many more!

Take some time to investigate what is available before starting.

# You want to write your own code?

## Directives

- OpenACC and OpenMP 4 define **directives** that can be used to instruct the compiler how to generate GPU code
- in theory the easiest path for porting

## GPU-specific Languages

- languages designed for GPU programming
- maximum flexibility and performance
- e.g. CUDA and OpenCL

# Things to consider

Before starting on a GPU implementation, it pays to ask some questions and do some preliminary exploration:

1. Is my program computationally or bandwidth intensive?
  2. Does it have enough parallel work to utilize the GPU?
  3. Must I change algorithms to expose enough parallelism?
  4. Are there serial bottlenecks that will limit scaling?
  5. Is the pain worth the gain?
- Questions 1, 2 and 3 will be discussed in this course.
  - Question 4 will be considered briefly here.
  - Question 5 requires answers for 1–4.

# Limitations to parallel speedup

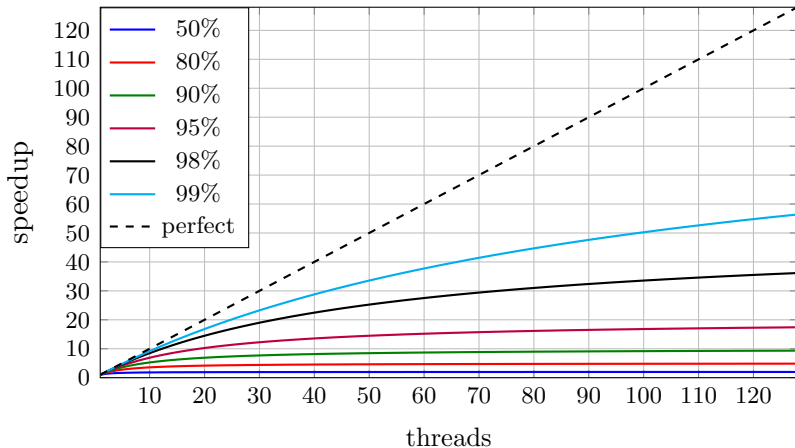
- Parallel speedup is limited by **the proportion of serial work** in your code.
- **Amdahl's law** defines the **maximum possible speedup** when only parts of the code can be parallelised

$$t_n = t_1 \left( p + \frac{(1-p)}{n} \right),$$

where  $t_n$  is time to solution for  $n$  threads and  $p \in [0, 1]$  is the proportion of sequential code.

- The limit on time to solution is  $\lim_{n \rightarrow \infty} = pt_1$ 
  - e.g. 1% of serial code gives a maximum 100× speedup.

# Amdahl illustrated





# CUDA

CUDA is a **parallel computing platform and API**

- For CUDA-enabled Nvidia GPUs.

We use CUDA as short hand for CUDA C/C++ and API

- CUDA C++ is a **superset** of C++
- Adds keywords for writing kernels to run on the GPU.
- Adds syntax for launching kernels on the GPU.

The CUDA toolkit is more than a programming language:

- Runtime API for managing GPU resources and execution.
- Tools including profilers and debuggers.

# Compiling CUDA

CUDA code is compiled with the **nvcc** compiler driver

- source files have `.cu` extension
- headers have `.h`, `.hpp`, `.hcu` extension.

CUDA compilation involves multiple splitting, compilation, preprocessing and merging steps

- `nvcc` hides this complexity from the user.
- It closely mimics the interface of the GNU compiler.
- Behind the scenes it:
  - uses GCC to compile the code that runs on CPU;
  - and compiles the GPU code separately.

# Compiling CUDA

## Example CUDA compilation

```
> nvcc -arch=sm_60 -lineinfo -O2 -std=c++11 -g -o foo foo.cu
```

Some flags are for **device** code generation:

- **-arch=sm\_60** target GPU architecture (Pascal)
- **-lineinfo** debug information for device code.

Some are for **host**:

- **-g** debug information for host code.

And some are for both **host and device**:

- **-O2** optimization level
- **-std=c++11** target language
- **-o foo** name of executable.

# Exercise: Getting Started on Piz Daint

In this exercise we will get introduced to Daint and make sure that everybody is set up.

```
# log on to daint with your course username & password
> ssh -X courseXX@daint

# get one node on the course reservation
<<<<<< HEAD
> salloc -C gpu --reservation=summer
=====
> salloc -C gpu --reservation=Pascal1day -p jalet
>>>>>> master

# go to scratch and get the course material
> cd $SCRATCH/my_repo
> git pull

# compile and test the demo
> cd topics/cuda/practicals/demos
> cat hello.cu
> module load cudatoolkit gcc
> nvcc -arch=sm_60 hello.cu -o hello
> srun ./hello
```